

TCSS 558 HW3 Design Document

Daniel McDonald, Jesse Carrigan and Sven Berger

December 2, 2013

Introduction

This project is meant to fulfill the basic requirements of the Chord protocol as outlined in the third and fourth homework literature. This document will explain the basic design of the entire project, and detailed design of those components necessary for this portion of the assignment.

Design Overview

Current Requirements

In addition to the requirements outlined in HW 3, in this assignment our chord protocol implementation must be extended to:

- Allow nodes to gracefully enter and exit the network.
- Allow a client to set data to a node in the network.
- Allow a client to retrieve data from a node in the network.
- Log the activity of every node in individual files.

The basic design from HW 3 was insufficient to accommodate the changes introduced in HW 4. Initially an attempt was made to closely follow the abstraction described in the reference Chord document, but due to the complexity of that implementation errors introduced went undetected, and after further development became intractably difficult to resolve. Once instability was introduced to the network, a bug was discovered in how the protocol determined the address range for each node. This problem went undetected due to the way the predecessor and successor were determined for the same node, masking the vulnerability. For these reasons much of the RMNode class required revision.

Adding and Removing Nodes From the Network

Nodes are added to the network in the same fashion as in HW3, but with the addition of logic to reassign the data in the range of the new node from the successor node formerly responsible for it.

When a node is removed, the node sends a signal to the predecessor indicating it should fix the finger table, and sends messages to all members of the local finger table to do the same. Next, it reassigns all the values in the key hash table to the successor node. It then sets itself to null. Although it does not remove itself from the RMI registry, any attempt to access it will result in a RemoteException, which is then caught and null is returned. Data will not be lost during this event.

Setting, Getting and Deleting Data on Nodes

A node receiving a Get request will first check to see if the hash of the requested key is within the range of virtual nodes it's responsible for. If it is, it returns the value from that

node. If it isn't, the node finds the successor for the given key, and passes the request on to it. Eventually, if the key is within the node space, the value will be found and returned. If it isn't, null will be returned.

When a node is asked to store a value at a specific key, much the same thing happens. Each node maintains an internal Map of all the virtual nodes within its range, and whatever values may be stored there. If the request is within this range, the value is stored at the hashed map value. If it is not, a similar process to a Get request occurs, and the data is passed to that node. If two keys should hash to the same value, whatever data is stored first will be overwritten by the second.

When a client requests data be deleted, the data is located the same way as it is for Get and Set, and the reference is removed from that node's internal map.

Logging Node Activity

When loaded, each node opens a file on the local machine, where the file name is the ID of the node. Whenever an action takes place for any node, it updates the file accordingly with a new record. A record consists of a timestamp, followed by relevant information. This file is never overwritten and holds the complete history of network activity.

An attempt was made to have nodes generate events observable by a special logging class, but this was found to be no better than the existing logging solution, and was ultimately abandoned.

Future Work

Homework 5 requires our chord implementation to be fault tolerant to the sudden loss of a node in the system. This means other nodes must re-learn their finger table automatically AND retain lost data. The requirements state that only one node at a time will ever be lost, so data redundancy could be as easy as storing the key values from either the predecessor node or successor, and then rebalancing this data when nodes join or leave.